

Word Embeddings

Giovanni Colavizza

Text Mining
Amsterdam University College

February 5, 2021

Announcements

- **TBD.**

Overview

1 Recap on Word Vectors

2 Word2Vec

3 Word2Vec Expanded

Recap on Word Vectors

Vector Semantics

“The meaning of a word is its use in the language.” Wittgenstein, 1953.

- ① *Pasta is best when cooked just right.*
- ② *Pizza should not be too cooked: it burns!*
- ③ Vector semantics combines two intuitions:
 - ▶ **Distributional approach**: define a word by the contexts it occurs into.
 - ▶ **Vectorize it**: use vectors to represent word meaning, as a point in space.
- ④ **Feature engineering** for NLP: word vectors are increasingly used as features for other tasks.
- ⑤ (Word) vectors are usually referred as (word) **embeddings** in modern neural network literature.

Co-occurrences

...ound and sonic power of a [new electric guitar played through] a guitar amp has play...
...[Some electric guitar models feature] piezoelectric pickups...
...[Playing guitar with a] pick produces a bright sound ...
...ings, he is known for [playing fretless guitar in his] performances...
...the neck of [a classical guitar is too] wide and the normal position ...
...t in the centre of Bristol [playing the piano , I was] punched in the head while, a...
...r in Houston, Texanstagram [playing the piano in his] flooded home after Hurrican H...
... some supplies, he stopped to [play the piano that was] sitting in knee-high water ...
...te and one black, who [played classical pianos together]...
...The [first electric from the] late 1920s used metal strin...
...technologies, for example [the electric and the] integration of mobile commun...
...study had each driver of [each electric drive unimpeded], perform a task whil...
...Honda to commence testing of [their new and the] American was no doubt more t...
...many design considerations for [the new car were "safety] innovations, performanc...
...would be possible if almost [all private cars requiring drivers], which are not in ...
... who donate to groups [providing private school scholarships have] written pieces att...
... that students participating [in private school choice programs] graduate high school...
...s in the establishment of this [new high school , named the] Gavirate Business School...
...Anna heads into her [final high school year before] university wanting somet...
... but he can prevent them from [playing at school]

Word-Context matrix


- We have a set of words V and a set of contexts they occur into C , taken from our corpus of documents. X in this case is a $|V| \times |C|$ matrix with word occurrences in contexts.
- The most intuitive context are co-occurrences with other words in V , within a certain **window**. In this case, X would be a $|V| \times |V|$ matrix.

	aardvark	...	computer	data	pinch	result	sugar	...
apricot	0	...	0	0	1	0	1	
pineapple	0	...	0	0	1	0	1	
digital	0	...	2	1	0	1	0	
information	0	...	1	6	0	4	0	

Figure 6.5 Co-occurrence vectors for four words, computed from the Brown corpus, showing only six of the dimensions (hand-picked for pedagogical purposes). The vector for the word *digital* is outlined in red. Note that a real vector would have vastly more dimensions and thus be much sparser.

Credit: J&M, ch. 6.

Vectors



Families of vectors

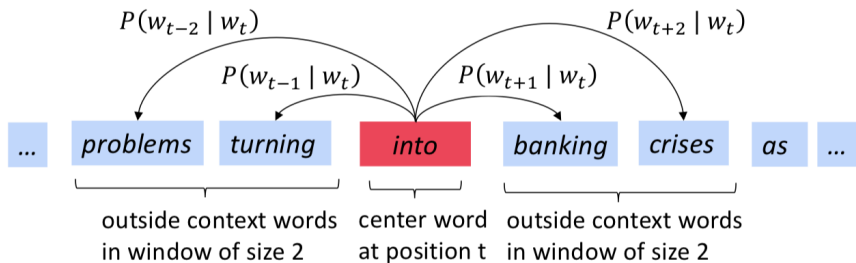
- **Sparse vectors:** many zero values and high-dimensional spaces. E.g., weighted co-occurrence matrices.
- **Dense vectors:** no zero values and comparatively smaller-dimensional spaces.
 - ▶ Dimensionality reduction (Singular Value Decomposition, Random indexing, Non-negative matrix factorization).
 - ▶ **Neural-network inspired (Word2Vec):** today.

Word2Vec

Families of vectors


- **Word2Vec**: a framework for learning dense word vectors.
- Idea:
 - ① We have a large corpus of text.
 - ② We want each word in the vocabulary to be represented by a vector.
 - ③ We can go through the corpus and establish a *context* o for every *center/focus word* c , using a certain window/span.
 - ④ **We use the similarity of the word vectors c and o to calculate the probability of context words o given c .**
 - ⑤ **We keep adjusting word vectors until our predictions are good.**

Words in context



Credit: Stanford CS224N.

Words in context



Credit: Stanford CS224N.

Words in context as data

Source Text

Training Samples

<div>The quick brown fox jumps over the lazy dog.</div>	→	(the, quick) (the, brown)
<div>The quick brown fox jumps over the lazy dog.</div>	→	(quick, the) (quick, brown) (quick, fox)
<div>The quick brown fox jumps over the lazy dog.</div>	→	(brown, the) (brown, quick) (brown, fox) (brown, jumps)
<div>The quick brown fox jumps over the lazy dog.</div>	→	(fox, quick) (fox, brown) (fox, jumps) (fox, over)

Credit: [http:](http://mccormickml.com/2016/04/19/word2vec-tutorial-the-skip-gram-model)

[//mccormickml.com/2016/04/19/word2vec-tutorial-the-skip-gram-model](http://mccormickml.com/2016/04/19/word2vec-tutorial-the-skip-gram-model).

The model

- Our task, for every c (center), o (context) pair, is to estimate high probabilities for:

$$p(w_o|w_c)$$

- *The model parameters are the word embeddings w .*
- For each word position $t = 1 \dots T$, we predict context words within a windows of size m , given the center word w_t (at each position):

$$L(\mathbf{w}) = \prod_{t=1}^T \prod_{-m \leq j \leq m; j \neq 0} p(\mathbf{w}_{t+j}|\mathbf{w}_t)$$

The model


- Loss function (of the negative log likelihood):

$$\mathcal{L}(\mathbf{w}) = -\frac{1}{T} \log L(\mathbf{w}) = -\frac{1}{T} \sum_{t=1}^T \sum_{-m \leq j \leq m; j \neq 0} \log p(\mathbf{w}_{t+j} | \mathbf{w}_t)$$

- *Minimizing the loss is equivalent to maximizing the likelihood.*
- How to calculate $p(\mathbf{w}_{t+j} | \mathbf{w}_t)$? Use two vectors for each word:
 - ▶ v_w when w is a center word
 - ▶ u_w when w is a context word
- Use the **Softmax** (generalization of the Sigmoid) to predict probabilities of a c (center), o (context) pair:

$$p(o|c) = \frac{\exp(u_o^T v_c)}{\sum_{w \in V} \exp(u_w^T v_c)}$$

Example



We learn to predict:

- $p(u_{problems} | v_{into})$
- $p(u_{turning} | v_{into})$
- $p(u_{banking} | v_{into})$
- $p(u_{crises} | v_{into})$
- ...

Credit: Stanford CS224N.

Softmax

Exponentiation makes anything positive

Dot product compares similarity of o and c .
 $u^T v = u \cdot v = \sum_{i=1}^n u_i v_i$
Larger dot product = larger probability

$$P(o|c) = \frac{\exp(u_o^T v_c)}{\sum_{w \in V} \exp(u_w^T v_c)}$$

Normalize over entire vocabulary to give probability distribution

- The Softmax maps any value to a probability distribution.
- It amplifies large values (*max*) but still gives non-zero probabilities to small values (*soft*).

Credit: Stanford CS224N.

Training via SGD

- Parameters: our word embeddings, **two per word**.
- Usually, these vectors have length d within 50-1000, thus $d \ll |V|$.
- Use gradient descent to optimize and find a minimum of the loss.




Figure 5.4 Visualization of the gradient vector in two dimensions w and b .

Credit: J&M, ch. 5.

Training via SGD

- Let us ignore for a moment the normalization term $\frac{1}{Z}$ and the external summations, which are straightforward.
- Let us take the first (partial) derivative w.r.t. v_c (similarly, you can do this for u_o):

$$\begin{aligned}\frac{\partial}{\partial v_c} \log \frac{\exp(u_o^T v_c)}{\sum_{w \in V} \exp(u_w^T v_c)} &= u_o - \sum_{x \in V} \frac{\exp(u_x^T v_c)}{\sum_{w \in V} \exp(u_w^T v_c)} \cdot u_x \\ &= u_o - \sum_{x \in V} p(x|c) \cdot u_x\end{aligned}$$

- Thus the derivative w.r.t. the central word vector v_c is the vector for the current context word u_o , minus the weighted average of the model's current representations of other possible contexts!
- *Derivation..*

Conclusion

- After having trained the model, we typically use the vectors v_w or the average of v_w and u_w .
- There are many good implementations of this model: *next lab*.
- This is a very rapidly advancing area, with many, more involved models now-a-days. Still, word embeddings are the main building block of deep learning NLP applications (used as **features**).
- Some more optional references and info below.

References

- Stanford CS224N classes 1 and 2:
<http://web.stanford.edu/class/cs224n/index.html>.
- Original Word2Vec paper
<https://arxiv.org/pdf/1301.3781.pdf>.
- Negative sampling paper <http://papers.nips.cc/paper/5021-distributed-representations-of-words-and-phrases-and-pdf>.
- Good tutorial <http://mccormickml.com/2016/04/19/word2vec-tutorial-the-skip-gram-model>.

Note: there is much more. Ask me if you are interested.

Word2Vec Expanded (optional)

Derivation for Softmax

- First, we need some notable derivatives:

$$\frac{\partial \log(x)}{\partial x} = \frac{1}{x}$$

$$\frac{\partial \exp(x)}{\partial x} = \exp(x)$$

$$\frac{\partial f(g(x))}{\partial x} = \frac{\partial f}{\partial g} \cdot \frac{\partial g}{\partial x} \rightarrow \text{chain rule}$$

Derivation for Softmax

- We can divide in two parts:

$$\frac{\partial}{\partial v_c} \log \frac{\exp(u_o^T v_c)}{\sum_{w \in V} \exp(u_w^T v_c)} = \frac{\partial}{\partial v_c} \log \exp(u_o^T v_c) - \frac{\partial}{\partial v_c} \log \sum_{w \in V} \exp(u_w^T v_c)$$

- First part:

$$\frac{\partial}{\partial v_c} \log \exp(u_o^T v_c) = u_o$$

- Second part:

$$\begin{aligned} \frac{\partial}{\partial v_c} \log \sum_{w \in V} \exp(u_w^T v_c) &= \frac{1}{\sum_{w \in V} \exp(u_w^T v_c)} \cdot \frac{\partial}{\partial v_c} \sum_{x \in V} \exp(u_x^T v_c) \\ &= \frac{\sum_{x \in V} \exp(u_x^T v_c) u_x}{\sum_{w \in V} \exp(u_w^T v_c)} \end{aligned}$$

Derivation for Softmax

- Combine:

$$\begin{aligned}\frac{\partial}{\partial v_c} \log \frac{\exp(u_o^T v_c)}{\sum_{w \in V} \exp(u_w^T v_c)} &= u_o - \frac{\sum_{x \in V} \exp(u_x^T v_c) u_x}{\sum_{w \in V} \exp(u_w^T v_c)} \\ &= u_o - \sum_{x \in V} p(x|c) u_x\end{aligned}$$

Negative sampling

Exponentiation makes anything positive

Dot product compares similarity of o and c .

$$u^T v = u \cdot v = \sum_{i=1}^n u_i v_i$$

Larger dot product = larger probability

$$P(o|c) = \frac{\exp(u_o^T v_c)}{\sum_{w \in V} \exp(u_w^T v_c)}$$

Normalize over entire vocabulary
to give probability distribution

- Normalizing over the entire vocabulary is very expensive.
- Idea: let us just create some **negative examples** (collocations absent in the data), and train a **binary logistic regression classifier** to distinguish between positive (real) and negative (fake) pairs.
- For every center-context pair, also sample K negative pairs. The center-context pair is going to be a positive datapoint, the negative pairs are negative datapoints.
- The logistic classifier then uses the same dot product of vectors as features, and a cross-entropy loss (see last class).

Variations of Word2Vec

- What we discussed is called the Continuous Bag Of Words model (**CBOW**).
- Alternatively, we can predict the center word using the context words: this is called the **Skip-gram** model.
- **GloVe**: approach that approximates global co-occurrence information instead (which we have seen in previous classes).

More references

- GloVe <https://nlp.stanford.edu/pubs/glove.pdf>.
- Word2Vec implicitly factorizes the (shifted) PPMI matrix we are familiar with: <https://papers.nips.cc/paper/5477-neural-word-embedding-as-implicit-matrix-factorization.pdf>.
- Evaluation of word embeddings:
<https://www.aclweb.org/anthology/D15-1036>.