

Creating annotated corpora

Giovanni Colavizza

Text Mining
Amsterdam University College

February 5, 2021

Announcements

- **TBD.**

Overview

1 NLP and Human Annotations

2 Evaluation


3 Platforms and shared tasks

NLP and Human Annotations

NLP and Human Annotations


- NLP (and ML in general) is driven by human-annotated corpora.
- <http://nlpprogress.com>.
- Annotation is **difficult** and **expensive**.

Annotation pipeline



Pustejovsky and Stubbs, 2012. Natural Language Annotation for Machine Learning.

Annotation pipeline



Pustejovsky and Stubbs, 2012. Natural Language Annotation for Machine Learning.

Specs and guidelines

- Goal: given our problem, *how can we formalize our description of the annotation process for multiple annotators to provide the same judgment?*
 - ▶ What is the goal of the project?
 - ▶ How will the annotation be created? (For example, which tags or documents to annotate first, how to use the annotation tools, etc.)
 - ▶ What is each tag called and how is it used? (provide examples and discuss problematic choices.)
 - ▶ What parts of the text do you want annotated?
- Note: annotation is usually boring and time-consuming, and cannot be done for 8 hours straight. Annotators also get better over time: early annotations might be discarded.

Pustejovsky and Stubbs, 2012. Natural Language Annotation for Machine Learning.

Adjudication

- **Adjudication** is the process of deciding on a single annotation for a piece of text, using information from all independent annotators.
- Yes, it is only possible when multiple annotators independently annotate (at least some) of the corpus. *This is a very good procedure to follow, and the only one which will allow to evaluate results.*
- It can be as time-consuming (or more so) as a primary annotation.
- It does not need to be identical with a primary annotation (all annotators can be wrong by chance), but unlikely so.

Pustejovsky and Stubbs, 2012. Natural Language Annotation for Machine Learning.

Evaluation

Interannotator agreement



annotator B

annotator A

	puppy	fried chicken
puppy	6	3
fried chicken	2	5

observed agreement = $11/16 = 68.75\%$

Credit: David Bamman, UCBerkeley.

Cohen's kappa

- Similar idea to mutual information: observed minus expected agreement.
- Cohen's kappa is defined for two annotators over the same set of annotation tasks:

$$\kappa = \frac{p_o - p_e}{1 - p_e}$$

Where p_o is the observed correct agreement and p_e the expected correct agreement.

		annotator A	
		puppy	fried chicken
annotator B	puppy	7	4
	fried chicken	8	81

Credit: David Bamman, UCBerkeley.

Cohen's kappa example

- $p_o = 0.88$
- $p_e = P(A = \text{puppy})P(B = \text{puppy}) + P(A = \text{chicken})P(B = \text{chicken})$
-

$$\kappa = \frac{0.88 - 0.773}{1 - 0.773} = 0.471$$

annotator A

		puppy	fried chicken
annotator B	puppy	7	4
	fried chicken	8	81

Credit: David Bamman, UC Berkeley.

Cohen's kappa scores

Note: these are rules of thumb.

0.80-1.00	Very good agreement
0.60-0.80	Good agreement
0.40-0.60	Moderate agreement
0.20-0.40	Fair agreement
< 0.20	Poor agreement

Exercise: try to calculate fringe cases. E.g., 50/50 puppet/chicken all in agreement, 0/100 puppet/chicken all in agreement, 50/50 wrong puppet/chicken all in agreement.

Credit: David Bamman, UC Berkeley.

Fleiss' kappa

- Extension to multiple annotators (> 2).
- Defined as Cohen's kappa but comparing pairs of annotators:

$$\kappa = \frac{P_o - P_e}{1 - P_e}$$

annotator A

	puppy	fried chicken
annotator B		
puppy	7	4
fried chicken	8	81

Credit: David Bamman, UC Berkeley.

Fleiss' kappa

- Number of annotators who assign category j to item i : n_{ij} .
- For item i with n annotations, how many annotators agree among all $n(n-1)$ possible pairs:

$$P_i = \frac{1}{n(n-1)} \sum_{j=1}^K n_{ij}(n_{ij} - 1)$$

- Note that N is the number of items, and K the available annotation categories. Average agreement among all items:

$$P_o = \frac{1}{N} \sum_{i=1}^N P_i$$

Fleiss' kappa, continued

- Probability of category j :

$$p_j = \frac{1}{Nn} \sum_{i=1}^N n_{ij}$$

- Expected agreement by chance:

$$P_e = \sum_{j=1}^K p_j^2$$

- Back to original formula:

$$\kappa = \frac{P_o - P_e}{1 - P_e}$$

Platforms and shared tasks

Annotation tools

- Brat <http://brat.nlplab.org>
- Inception <https://inception-project.github.io>
- Prodigy <https://prodi.gy>

Annotation platforms

- Supervisely <https://supervise.ly>
- Dataturks <https://dataturks.com>
- Amazon Mechanical Turk <https://www.mturk.com>
- Figure Eight <https://www.figure-eight.com>